

# Flink in Genomics

Efficient and scalable processing of raw Illumina BCL data

F. VERSACI   L. PIREDDU   G. ZANETTI

– FlinkForward 2016 –

13 September 2016



- 1 Introduction
- 2 BCL to FASTQ Conversion
- 3 Implementation in Flink
- 4 Evaluation and Final Considerations




- Research center in Sardinia, Italy
- Focus on **big data**, **biosciences**, HPC, visual computing, energy and environment

# Next-Generation Sequencing

## Cost

- Genome sequencing is now **much cheaper** than in the past
- About **1000 euros** per **whole human genome**



(Data from <https://www.genome.gov/sequencingcosts/>)

# Next-Generation Sequencing

## Applications

High-throughput DNA sequencing has many applications, including

- Research into understanding human genetic diseases
- Medicine, e.g., oncology, clinical pathology, ...
- Human phylogeny
- Personalized diagnostic applications

### Huge amount of data

A single sequencer can produce 1 TB/day of data

- Which need to be converted, filtered, aggregated, reconstructed, analysed, ...

# Standard pipeline

When using **Illumina sequencers**, the standard pipeline starts with two programs:

**bcl2fastq2** Proprietary, open-source tool by Illumina to **convert raw BCL data** to FASTQ format

**BWA-MEM** Free (GPLv3) **aligner** to reconstruct the full genomic sequence based on the **short reads** generated by the sequencer

## Problem

- Parallel tools, but **shared-memory** (single node)
- To exploit **more nodes** data need to be **distributed**, there can be **failures**, etc.

In this talk we present a **distributed-memory** BCL converter

- Developed within the **Flink framework**
- Written in **Scala**
- **Efficient** (i.e., speed comparable to bcl2fastq2)
- **Scalable**
- Can easily be **integrated** into existent **Hadoop/YARN workflows**



# Outline

- 1 Introduction
- 2 BCL to FASTQ Conversion
- 3 Implementation in Flink
- 4 Evaluation and Final Considerations



# Shotgun genome sequencing

- The DNA is a sequence of **four bases**: Adenine, Cytosine, Guanine and Thymine (**A, C, G and T**)
- To **reconstruct** it, the genome is broken up into **short fragments (reads)**
- The fragments are **attached** to a support (**tile**)
- **Fluorescent molecules** are iteratively attached to bases of the DNA fragments being sequenced
- At **each cycle**, the machine acquires (optically) a single base from all the fragments




# File organization

- We adopt the file structure of **Illumina HiSeq 3000/4000** machines
- A single file refers to data obtained by specific **lane**, **tile** and **cycle** combination
- E.g., file `L003/C80.1/s_3_1213.bc1.gz` corresponds to data read from **tile  $t = 1213$** , in **lane  $l = 3$**  during **cycle  $c = 80$**
- **Test dataset**:  $8 \text{ lanes} \times 112 \text{ tiles/lane} \times 210 \text{ cycles} = 188,160$  gzip-compressed BCL files = about **250 GB**

# BCL to FASTQ conversion

- BCL files are **arrays of bytes**
- Each byte encodes a **base (bits 0-1)** and a **quality score (bits 2-7)**
- **.filter** files specify which reads should be ignored
- **.locs** files contain some **metadata** which need to be attached to each read
- To get the reads from the raw BCL data we need to perform some sort of **matrix transposition**



# Outline

- 1 Introduction
- 2 BCL to FASTQ Conversion
- 3 Implementation in Flink**
- 4 Evaluation and Final Considerations

## Implementation choices

- The converter is written in **Scala**
- We use **sbt** to handle **compilation** and **dependencies**
- All the code is less than **1000 lines**
- No fancy IDEs, just **EMACS** as editor

[illegible]

# Algorithmic overview

For each lane/tile combination

- BCL files corresponding to different cycles are opened concurrently
- Bases and quality scores are extracted and filtered
- For each fragment a text header is added, containing various meta-data and an index
- Then they are sorted by their indexes
- Since there can be read errors also in the indices, the repartition is fuzzy: a parameter sets the numbers of allowed misinterpreted symbols
- Finally, a gzip-compressed file for each index is written to disk

# DataSet vs DataStream

- Our data is **static** and read from a **storage** unit
- This is **not** a typical **streaming** application
- We have tried both **DataSet** and **DataStream** structures
- Using **DataStream** is faster
- Because of its better **overlap of I/O and computations?**

## Lesson Learned #1

**Try DataStream** even if it doesn't seem like a natural fit for your application

# Data granularity

- BCL files are **arrays of bytes**
- It might seem natural to process them in Flink as **DataStream[Byte]**
- But reading and writing single bytes is **not efficient**
- We process data in **bigger chunks** (2048 bytes)
- It imposes a **lower load** on the streaming framework
- Better **cache locality** exploitation

## Lesson Learned #2

Avoid fine granularity and read in **larger chunks**



# Job granularity

- The **job unit (mini-job)** is the processing of a lane-tile combination
- Mini-jobs run for about **one minute** on one core
- We can choose to aggregate **n mini-jobs** into a Flink job
- And assign **c cores** to each Flink job
- E.g., aggregate **n = 16** mini-jobs and run them on **c = 4** cores
- What about launching **one huge Flink** job which handles all the work and cores?
- Best results with **n = 2** and **c = 1** (with processor **SMT = 2**)

## Lesson Learned #3

Keep Flink jobs **reasonably small**

# Low level optimizations

## Use of ByteBuffer

- To extract bases and quality scores we need to perform some **bit masks and shifts**
- E.g., to get the quality score from **byte b** we can run

```
val b : Byte = in.get
val q : Byte =
    (0x21 + (b & 0xFC) >>> 2).toByte
```

- We can obtain a **8x speed-up** by grouping bytes into **64-bit longs** and executing the equivalent operations:

```
val r : Long = in.getLong
val q : Long = 0x2121212121212121L
    + ((r & 0xFCFCFCFCFCFCFCFCFL) >>> 2))
```

- To interpret **byte arrays as longs**, we need to use the **ByteBuffer class**

# Low level optimizations

## Use of look-up tables

- We need to convert bases from numeric to ASCII notation
- E.g., 0x0001020303020100 maps to "ACGTTGCA"
- We can do it efficiently by compressing the input and using it as an index in a look-up table
- E.g., 0x0001020303020100 is compressed to index 0b0001101111100100 = 0x1BE4 and searched in the precomputed look-up table
- The table has  $2^{16} = 65536$  entries

- To **schedule** the Flink jobs we use Scala **Futures**
- **Parallel** and **not blocking**

```
// numTasks = number of Flink jobs  
implicit val ec = scala.concurrent.ExecutionContext  
    .fromExecutor(Executors.newFixedThreadPool(numTasks))  
val miniJobs : Seq[MiniJob] = reader.getMiniJobs  
// fpar = number of mini-jobs per Flink jobs  
val flinkJobs = work.sliding(fpar, fpar)  
    .map(fj => Future{runJobs(fj)})  
// convert Seq[Future] to Future[Seq]  
val flist = Future.sequence(flinkJobs)  
scala.concurrent.Await.result(flist, Duration.Inf)
```

# Outline

- 1 Introduction
- 2 BCL to FASTQ Conversion
- 3 Implementation in Flink
- 4 Evaluation and Final Considerations

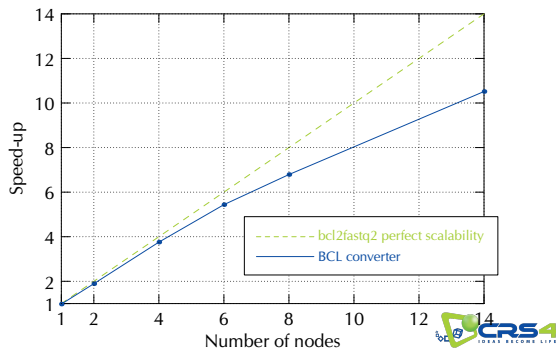
- Experiments run on the **Amazon Elastic Compute Cloud (EC2)**
- Up to **14** instances of **r3.8xlarge** machines
  - CPU**s 32 virtual cores (Xeon E5-2670 v2, 25 MB cache)
  - RAM** 250 GB
  - Disks** 2x320 GB SSD
  - Network** 10 Gb Ethernet
- HDFS distributed among the **n computing nodes**
- Each datanode using its **two SSD disks**
- **YARN** running on the same **n nodes**
- **Flink** running inside **Hadoop/YARN**

# Results

Strong, absolute scalability (bcl2fastq2 as baseline)

- Running time of **Illumina bcl2fastq2**: 57.1 minutes on a single node
- bcl2fastq2 is written in **C++** and exploits **Boost libraries**
- $896 = 64 \cdot 14$  total Flink jobs

Nodes	Time (minutes)
1	58.4
2	30.2
4	15.2
6	10.5
8	8.4
14	5.4



# I/O vs Computations

- The program is **CPU-bound** on the tested hardware
- Total **I/O size** (input+output):  $\approx 500$  GB
- I/O rate on **single node**:  $\approx 150$  MB/s
- I/O rate on **14 nodes**:  $\approx 1.6$  GB/s
- Note: both input and output are **gzip-compressed**



# Flink features – What we have used

## Custom Flink Input/OutputFormat

**Hadoop libraries** to read/write files

```
import org.apache.hadoop.fs.  
    {FileSystem, FSDataInputStream, FSDataOutputStream, Path}  
import org.apache.hadoop.io.compress.  
    {CompressionCodecFactory, CompressionInputStream}  
import org.apache.hadoop.io.compress.zlib.  
    {ZlibCompressor, ZlibFactory}
```

## DataStream

- map and flatMap
- MapFunction and FlatMapFunction
- filter
- split and select

# Flink features – What was not available

## Efficient zip of DataStreams

### Problem

- Given two data streams

```
val names: DataStream[String]
```

```
val ages: DataStream[Int]
```

- Join them as

```
val combined: DataStream[(String, Int)]
```

- Useful when reading data about the **same object from different files**
- E.g., .bcl, .locs and .filter files
- Inverse** function of

```
val names = combined.map(_._1)
```

```
val ages = combined.map(_._2)
```

# Flink features – What was not available

A smarter job scheduler

## Remark

We're talking about **Flink 1.0**: it seems the new job scheduler is much smarter :)

It would be convenient for the **job scheduler** to be able to

- Pick jobs from some (priority?) **queue**
- Runs them **concurrently** on the available Flink task slots
- Start a new job as soon as another one finishes
- Handle **failures** and **retries**

- Integrate our converter into **Seal**<sup>1</sup> toolkit for short DNA reads manipulation and analysis
- Adopt Flink also in the **second stage** of the pipeline, i.e., have a Flink-based **aligner**

---

<sup>1</sup><http://biidoop-seal.sourceforge.net/>

- Integrate our converter into **Seal**<sup>1</sup> toolkit for short DNA reads manipulation and analysis
- Adopt Flink also in the **second stage** of the pipeline, i.e., have a Flink-based **aligner**

Thanks for your attention!

---

<sup>1</sup><http://biidoop-seal.sourceforge.net/>