

Kung

RBea: Scalable Real-Time Analytics at King

Gyula Fóra
Data Warehouse Engineer
Apache Flink PMC



About King

We make awesome mobile games

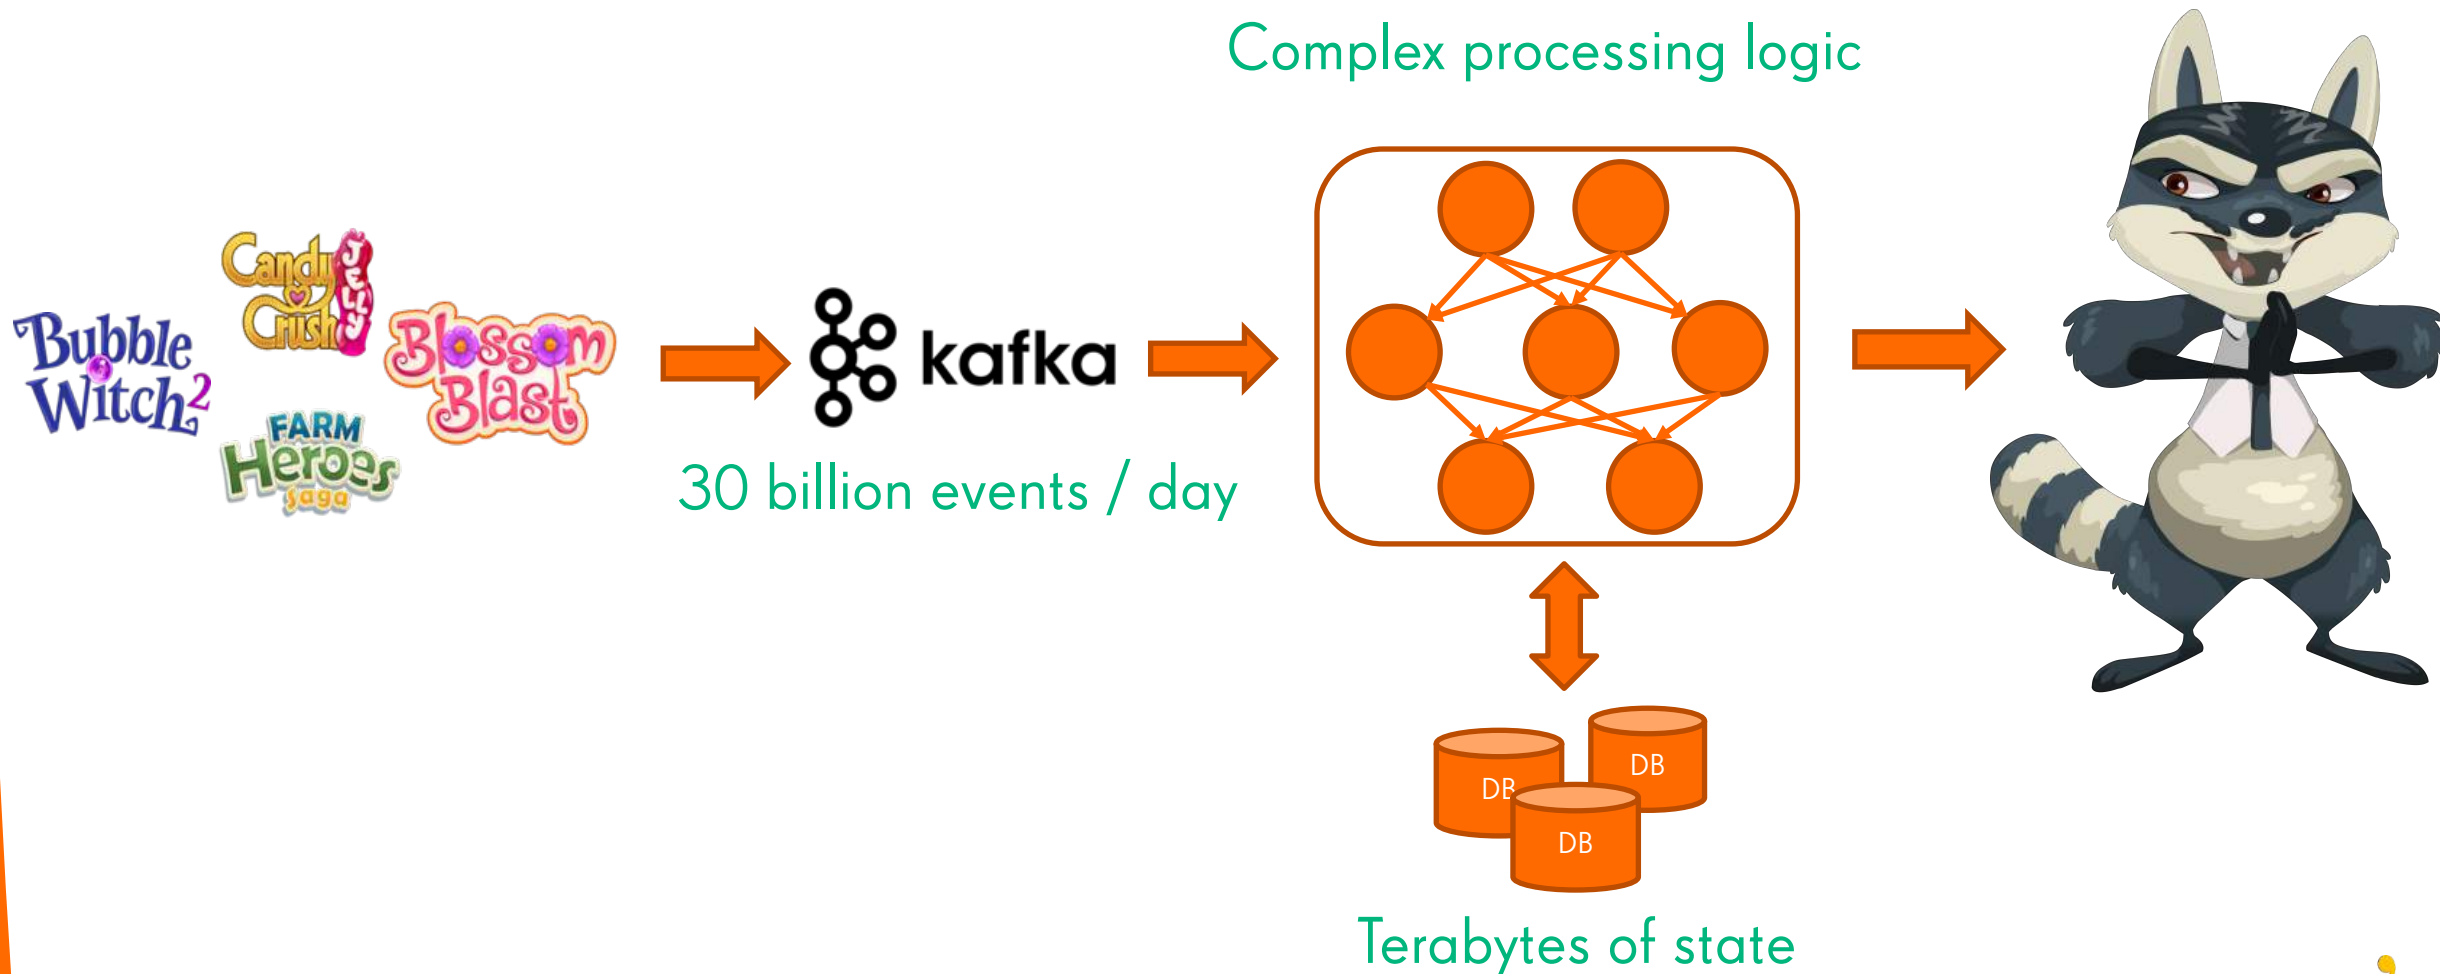
463 million monthly active users

30 billion events per day

And a lot of data...



From streaming perspective...



How do we use Flink



Standalone deployment

Few heavy streaming jobs

RocksDB state backend with caching

Lot of custom tooling

The RBea platform

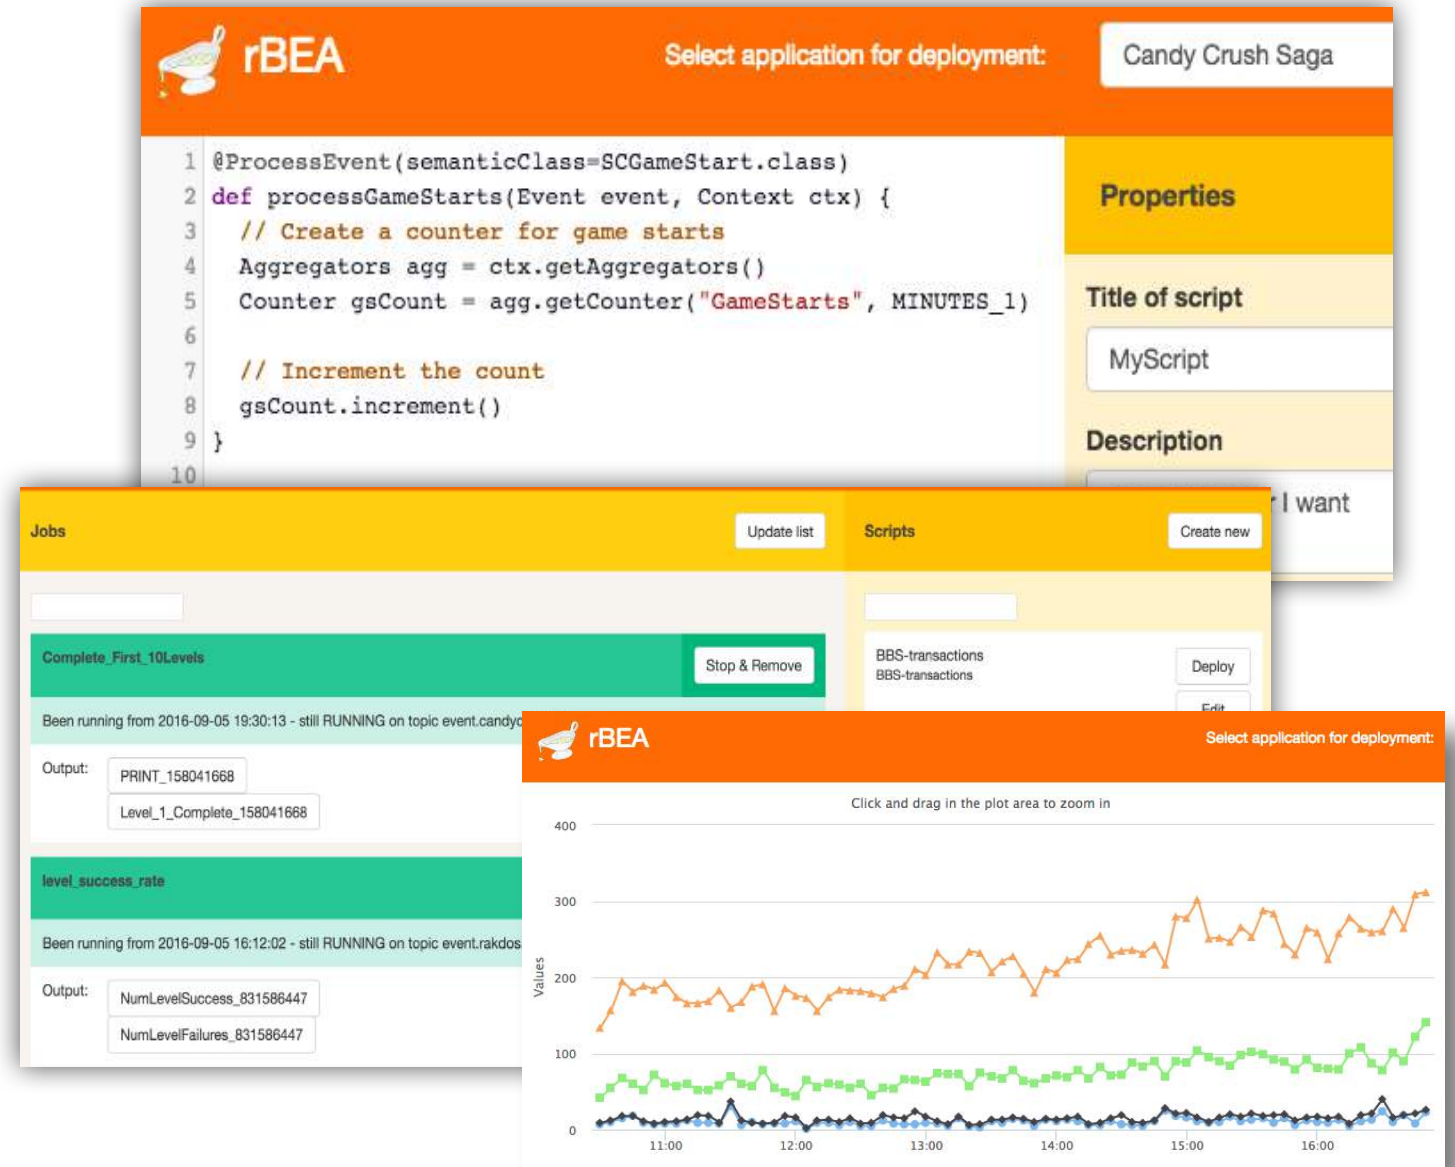
Scripting on the live streams

Deploy from browser/python

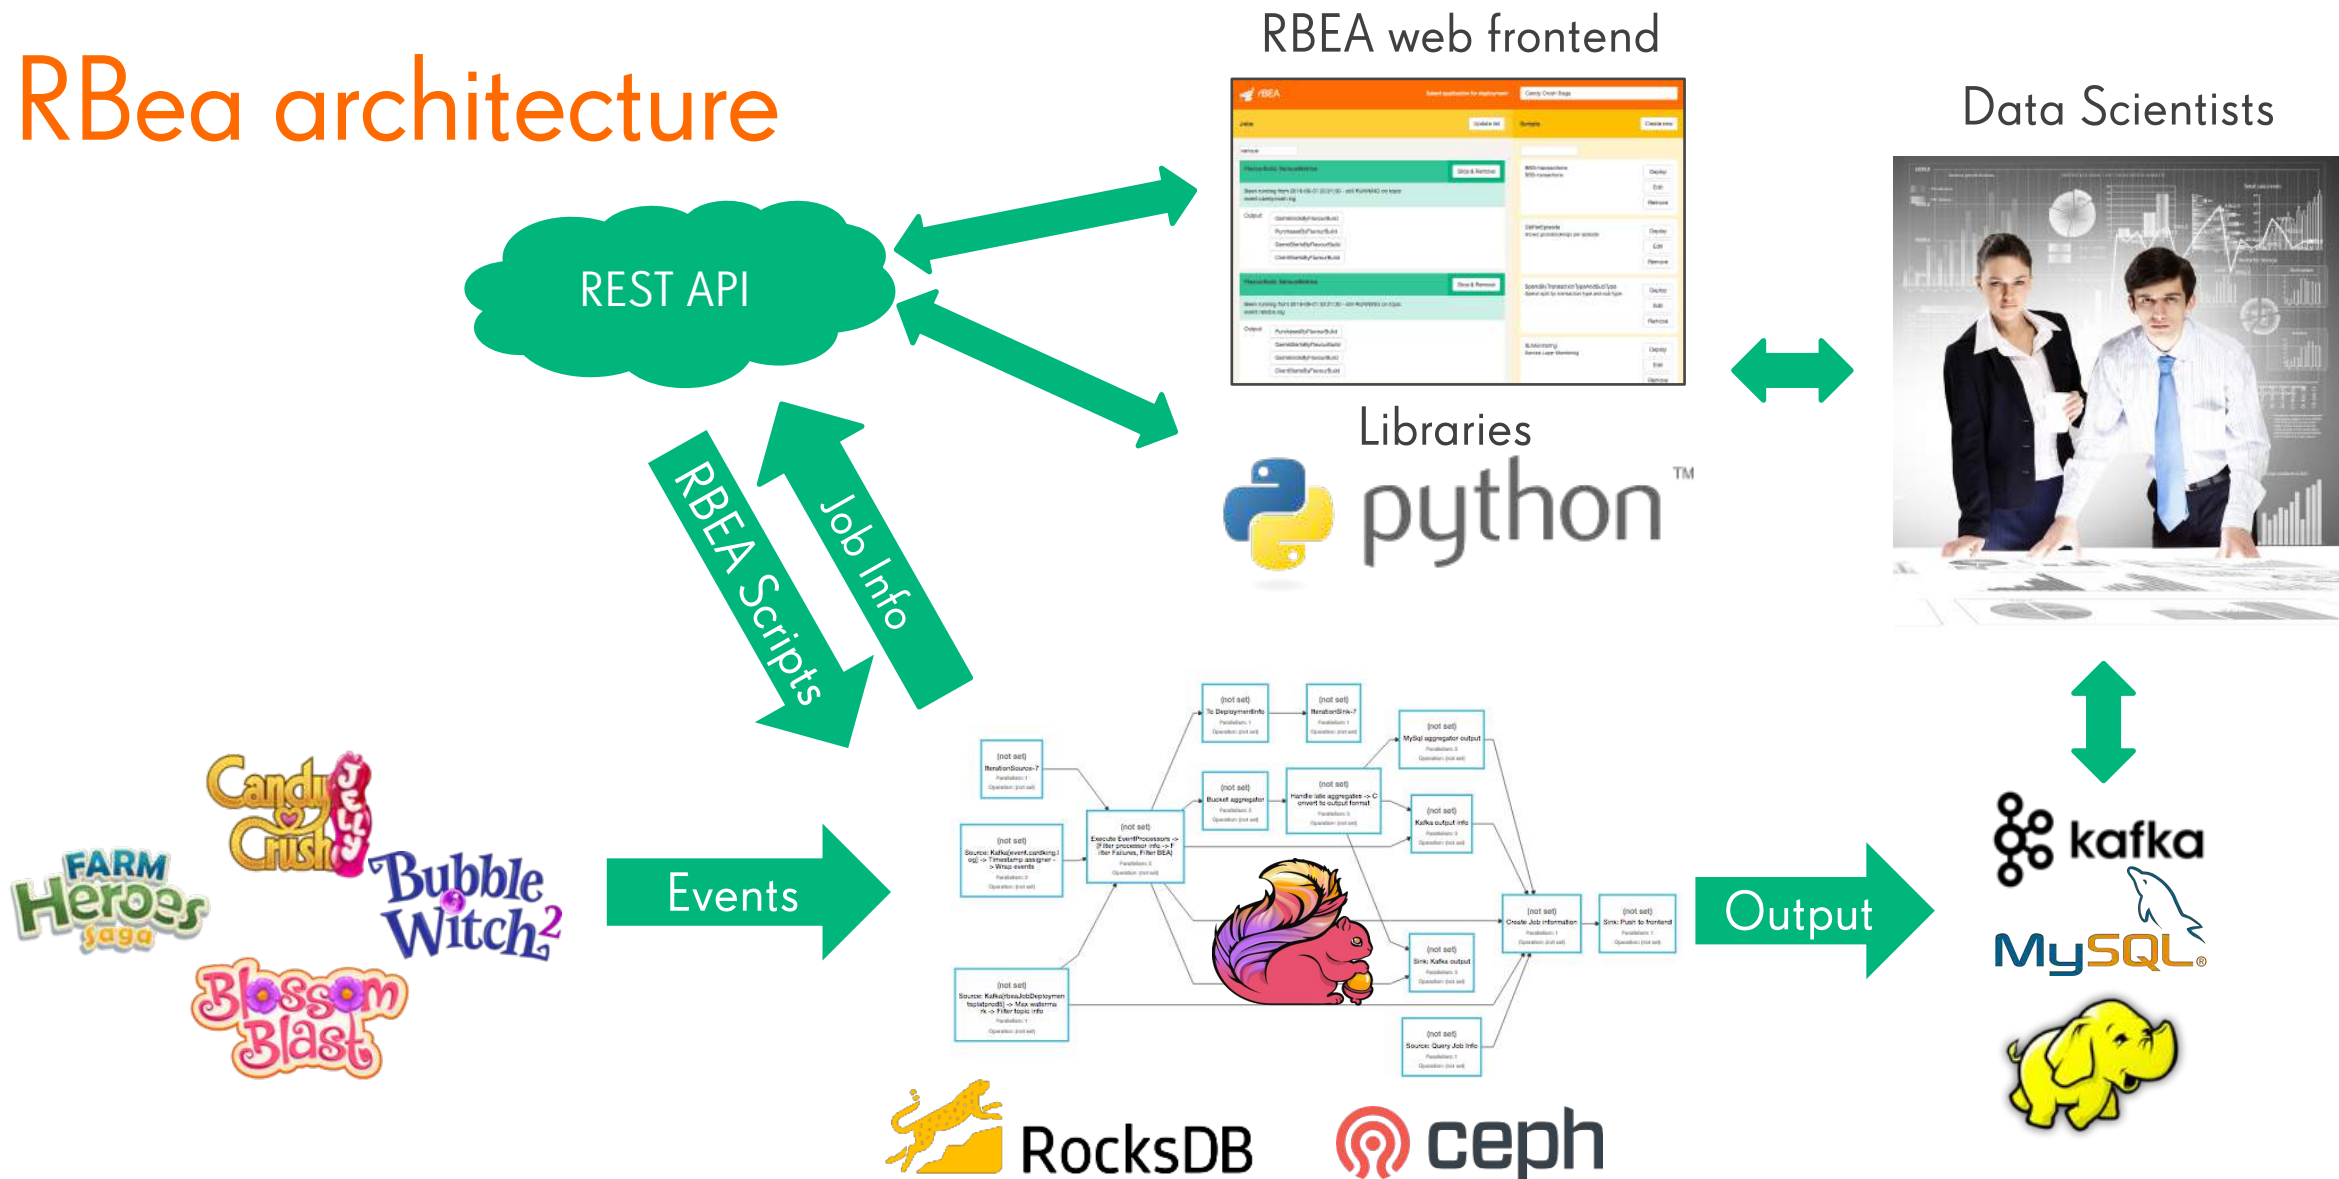
Window aggregates

Stateful computations

Scalable + fault tolerant



RBea architecture



RBea backend implementation

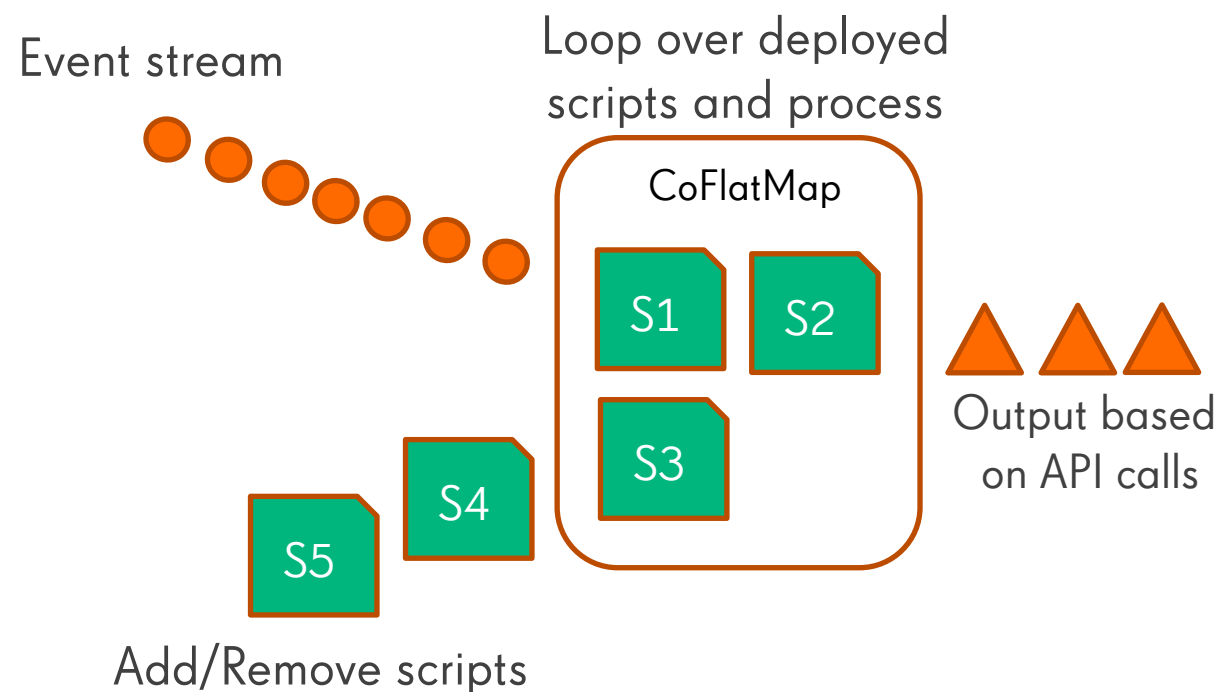
One stateful Flink job / game

Stream events and scripts

Events are partitioned by user id

Scripts are broadcasted

Output/Aggregation happens downstream



Dissecting the DSL

```
@ProcessEvent(semanticClass=SCPurchase.class)
def process(SCPurchase purchase,
            Output out,
            Aggregators agg) {

    long amount = purchase.getAmount()
    String curr = purchase.getCurrency()
    out.writeToKafka("purchases", curr + "\t" + amount)

    Counter numPurchases = agg.getCounter("PurchaseCount", MINUTES_10)
    numPurchases.increment()
}
```



Dissecting the DSL

```
@ProcessEvent(semanticClass=SCPurchase.class)
def process(SCPurchase purchase,
            Output out,
            Aggregators agg) {
```

```
    long amount = purchase.getAmount()
    String curr = purchase.getCurrency()
    out.writeToKafka("purchases", curr + "\t" + amount)
```

```
    Counter numPurchases = agg.getCounter("PurchaseCount", MINUTES_10)
    numPurchases.increment()
}
```

Processing methods by annotation

Event filter conditions

Flexible argument list

Code-generate Java classes

=> void processEvent(Event e, Context ctx);



Dissecting the DSL

```
@ProcessEvent(semanticClass=SCPurchase.class)
```

```
def process(SCPurchase purchase,
```

```
    Output out,
```

```
    Aggregators agg) {
```

```
    long amount = purchase.getAmount()
```

```
    String curr = purchase.getCurrency()
```

```
    out.writeToKafka("purchases", curr + "\t" + amount)
```

```
    Counter numPurchases = agg.getCounter("PurchaseCount", MINUTES_10)
```

```
    numPurchases.increment()
```

```
}
```

Output calls create *Output* events

Output(KAFKA, "purchases", "...")

These events are filtered downstream and sent to a Kafka sink



Dissecting the DSL

```
@ProcessEvent(semanticClass=SCPurchase.class)
```

```
def process(SCPurchase purchase,
```

```
    Output out,
```

```
    Aggregators agg) {
```

```
    long amount = purchase.getAmount()
```

```
    String curr = purchase.getCurrency()
```

```
    out.writeToKafka("purchases", curr + "\t" + amount)
```

```
        Counter numPurchases = agg.getCounter("PurchaseCount", MINUTES_10)
```


```
        numPurchases.increment()
```

```
}
```

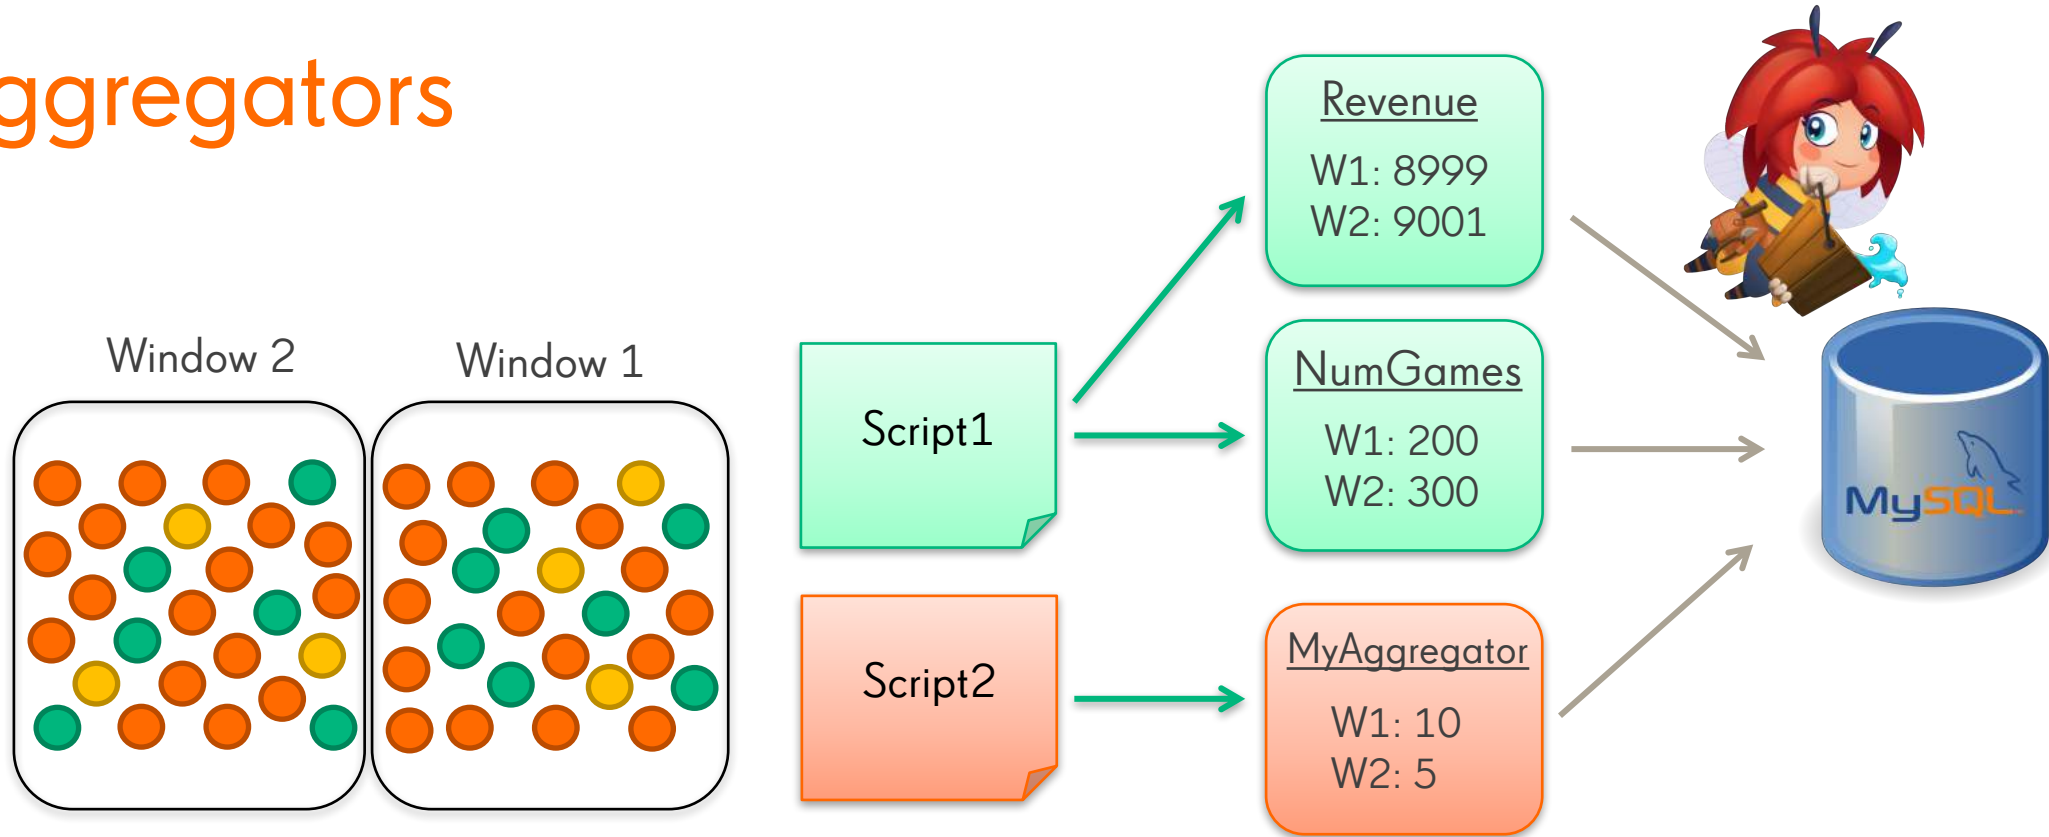
Aggregator calls create *Aggregate* events

Aggr (MYSQL, 60000, "PurchaseCount", 1)

Flink window operators do the aggregation



Aggregators

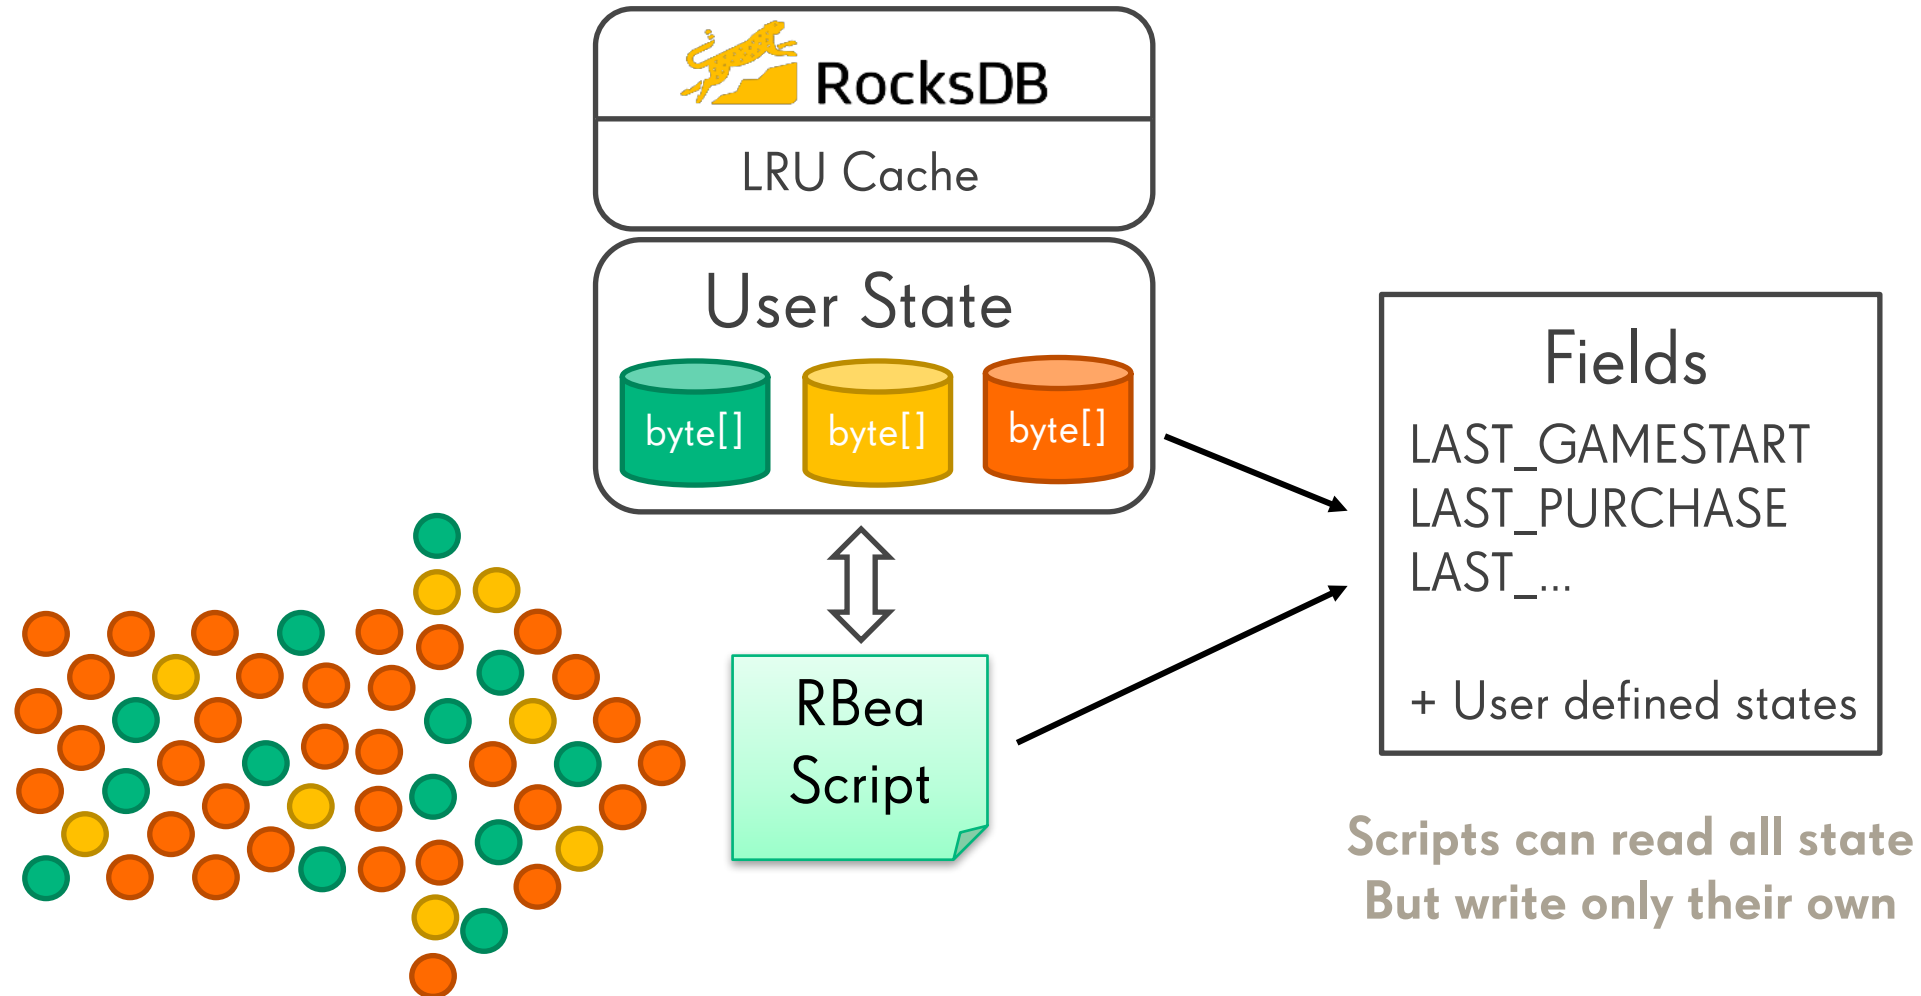


Event time windows
Window size / aggregator

Dynamic window assignment

```
long size = aggregate.getWindowSize();  
long start = timestamp - (timestamp % size);  
long end = start + size;  
TimeWindow tw = new TimeWindow(start, end);
```

User states



Things you might wonder...

Can slow scripts affect other scripts?

Yes, but we are working on it

Separate test/live environments

What does the backend know about the scripts?

Outputs produced

Failures (causes) => these are propagated

Runtime stats in the future

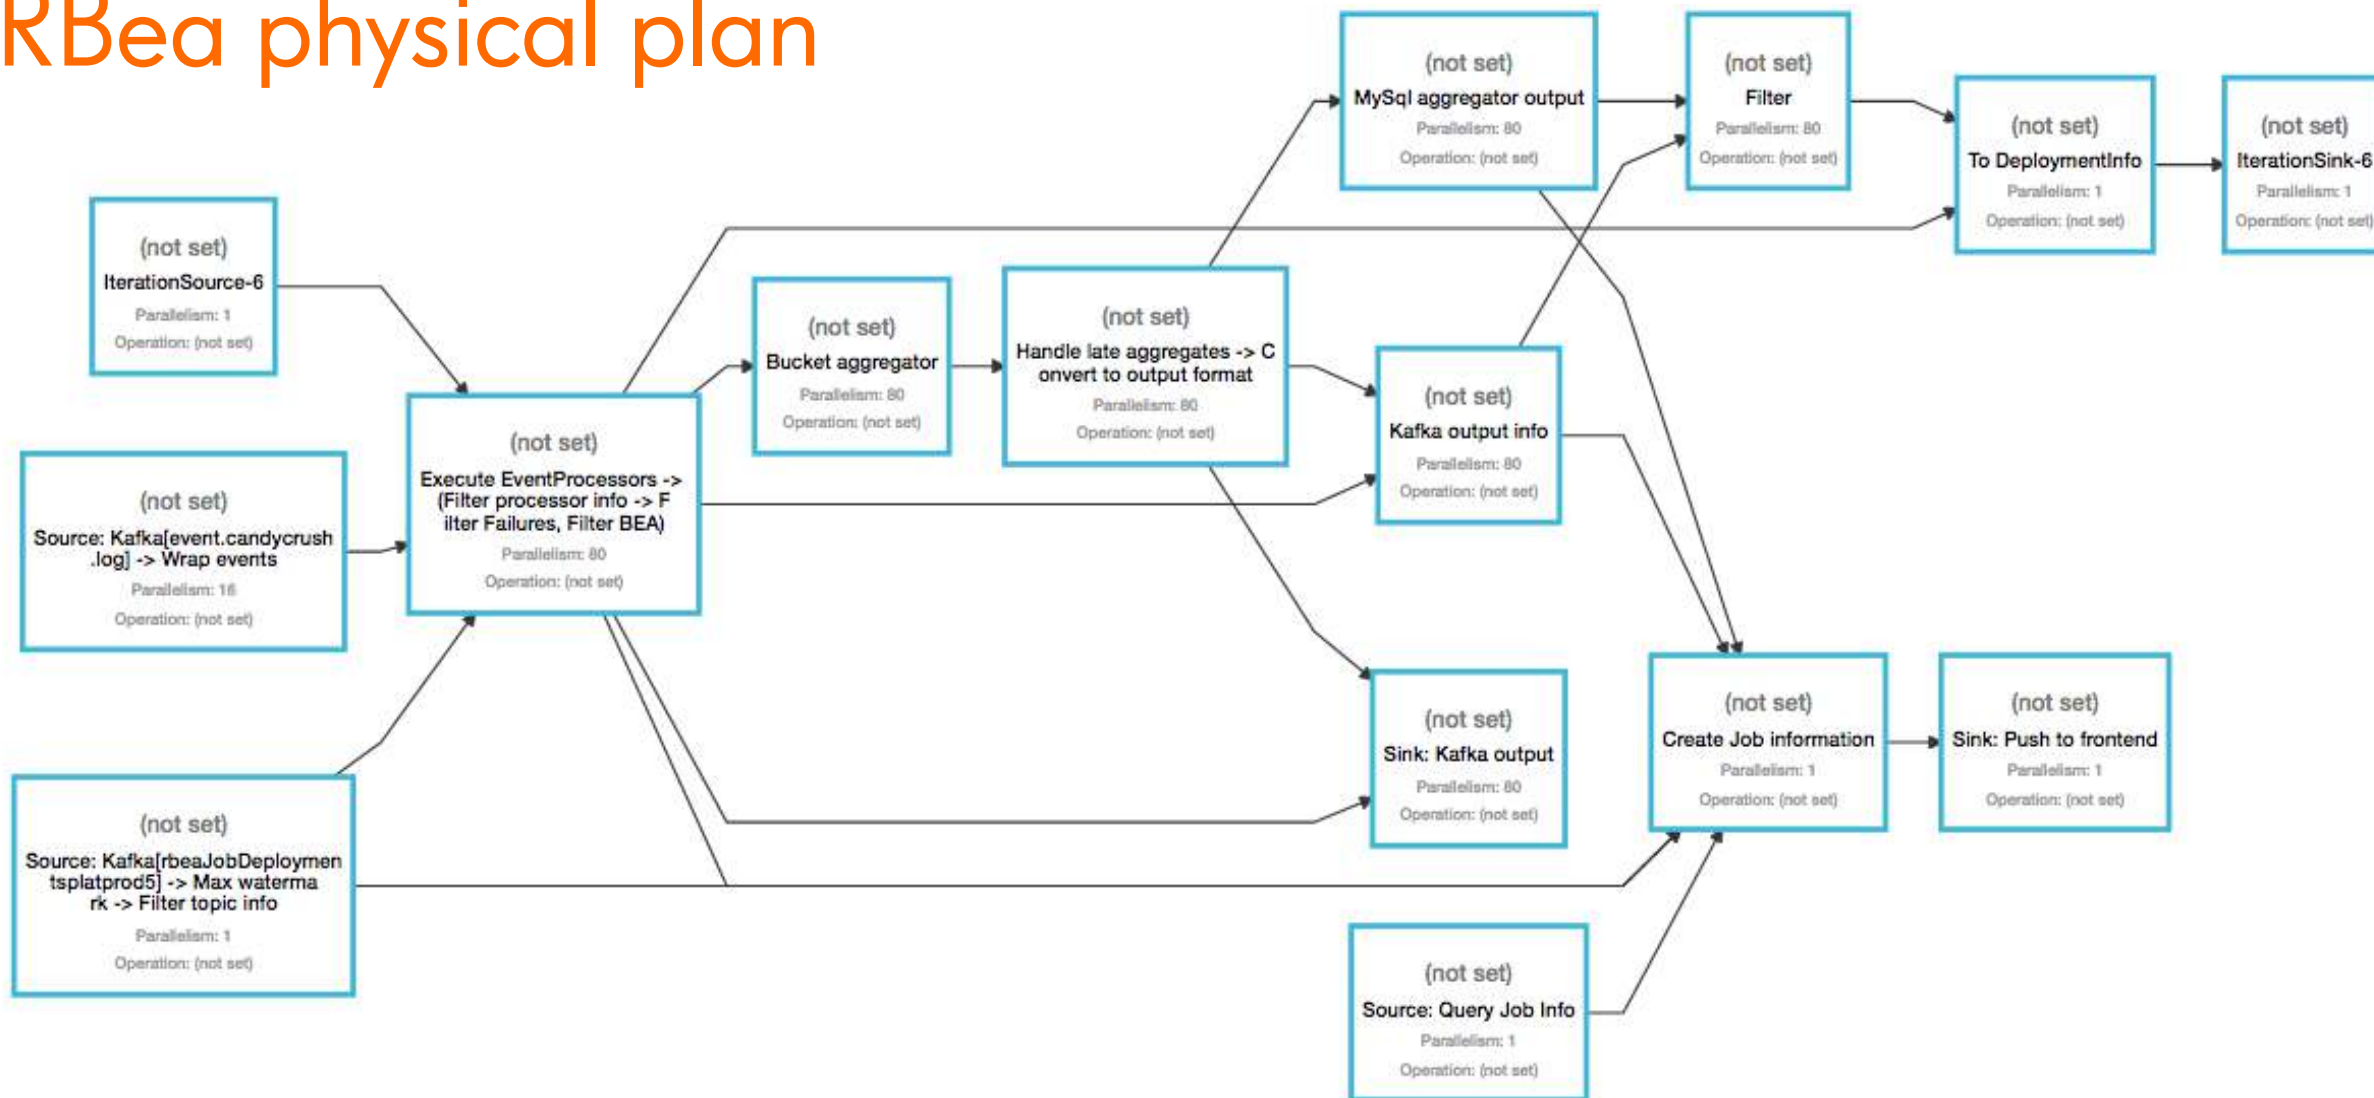
Is RBea useful compared to custom Flink jobs?

Writing + maintaining streaming jobs is hard

Especially with state, windowing, MySQL etc.



RBea physical plan



Wrap up

RBea makes streaming accessible to every data scientist at King

We leverage Flink's stateful and windowed processing capabilities

People love it because it's simple and powerful



Thank you!